

BATTLE OF PONG: DQN VS PG

Yash, Matt, Kenny, John

PROBLEM BACKGROUND


Given

1. Pixels from a game of pong
2. An indication of wins and losses
3. An opponent agent

Can we create an agent that could beat a human player?

PROBLEM MOTIVATION

- PONG is a classic game with an early, simple AI
- Games are an easy way to test theories
- Wanted to find the BEST training algorithm for this game



GOAL

To see what performs better in a basic game of pong

1. Deep Q-Learning Network (DQN)
 - a. Deep Q-Learning (DQN): “Human-level control through deep reinforcement learning”. 2015.
 - i. DeepMind, AlphaGo
2. Policy Gradient (PG)
 - a. Policy Gradient (PG): “Asynchronous Methods for Deep Reinforcement Learning”. 2016

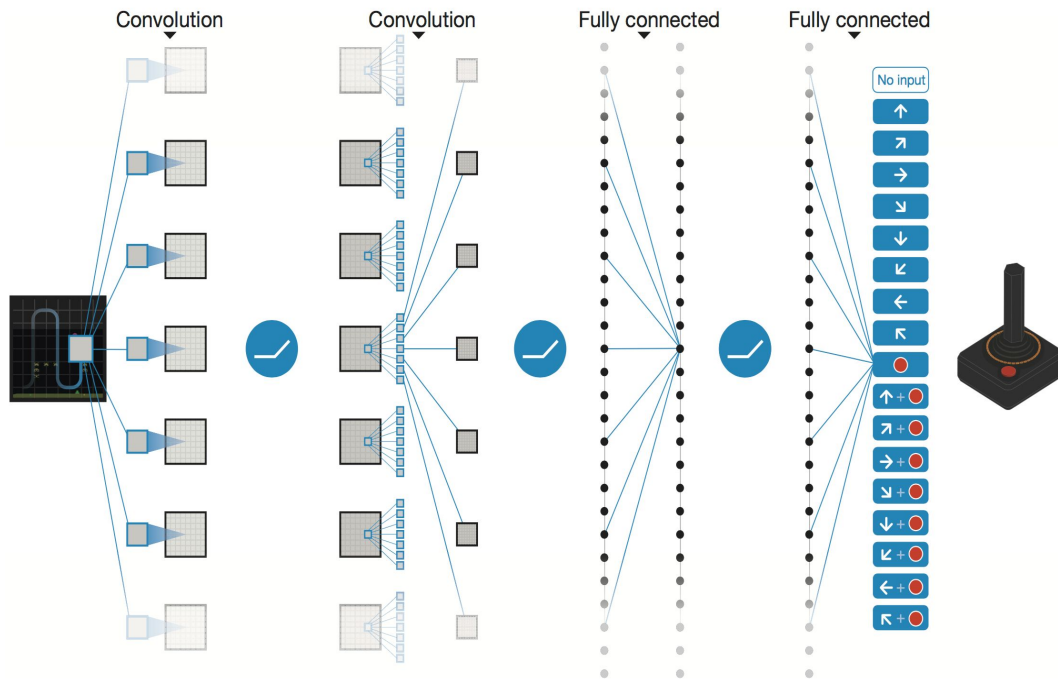
ALGORITHM: DEEP Q-LEARNING NETWORK

- Choose an action based on the best Q-value available
- Network architecture: Convolutional Neural Network to process images

ARCHITECTURE: DEEP Q-LEARNING NETWORK


- **Convolutional neural network:**

- Overlapping nodes
- Captures visual input in small chunks
- Final output comes from a regular neural net



ARCHITECTURE: DEEP Q-LEARNING NETWORK

- Chunking the input image and look for patterns
- Match sub-sections (features) against known outputs in similar spatial regions
- Does this by trying to match the feature at every possible location
- For PONG, we look only for features we care about, such as the paddles and the ball



ALGORITHM: DEEP Q-LEARNING NETWORK

- Explore vs Exploit
 - ϵ - choosing a random action
- Execute the action



ALGORITHM: DEEP Q-LEARNING NETWORK

- Observe & record in replay memory
 - State
 - Action
 - Reward
 - New State

ALGORITHM: DEEP Q-LEARNING NETWORK

- Sample from replay memory
 - Back propagate error
 - Reduces correlations between updates
- Repeat process
 - Unless game has ended




ALGORITHM: DEEP Q-LEARNING NETWORK

- Previous Uses
 - DeepMind - <https://deepmind.com/research/dqn/>



POLICY GRADIENT ARCHITECTURE

Adopted from Dr. Andrej [Karpathy](#)



- Take in images from the game and preprocess them.
- Use the Neural Network to compute a probability of moving up.
- Sample from that probability distribution and tell the agent to move up/down.
- If the round is over, find whether you won or lost.
- When the episode has finished (someone got to 21 points), pass the result through the backpropagation algorithm to compute gradient for the weights.
- After 10 episodes have finished, sum up the gradient and move the weights in the direction of the gradient.
- Repeat this process until weights are tuned enough.

HYPERPARAMETERS/INITIALIZATIONS


```
batch_size = 10 # every how many episodes to do a param update?
gamma = 0.99 # discount factor for reward (i.e later rewards are exponentially less important)
decay_rate = 0.99 # decay factor for RMSProp leaky sum of grad^2
num_hidden_layer_neurons = 200 # number of hidden layer neurons
learning_rate = 1e-4 # for convergence (too low- slow to converge, too high, never converge)
resume = True # resume from previous checkpoint?
render = False
```

- Weight initialization

```
input_dimensions = 80 * 80 # input dimensionality: 80x80 grid (the pong world)
weights = {} # initialize weights
# By using Xavier initialization, we make sure that the weights are not too small but not too big to
# propagate accurately the signals.
weights['W1'] = np.random.randn(num_hidden_layer_neurons, input_dimensions) / np.sqrt(
    input_dimensions) # "Xavier" initialization
weights['W2'] = np.random.randn(num_hidden_layer_neurons) / np.sqrt(num_hidden_layer_neurons)
```

FORWARD PASS THROUGH THE NEURAL NET

- Compute `hidden_layer_values` to get initial activation values.
- `ReLU(hidden_layer_values)`.
- Compute `output_layer_value` to get the probability of going up.
- `p = sigmoid(output_layer_value)` to make sure the probability is between 0 and 1.

```
def apply_neural_nets(observation_matrix, weights):  
    """ Based on the observation_matrix and weights,  
    compute the new hidden layer values and the new output layer values. """  
    hidden_layer_values = np.dot(weights['W1'], observation_matrix)  
    hidden_layer_values = relu(hidden_layer_values)  
    output_layer_values = np.dot(hidden_layer_values, weights['W2'])  
    output_layer_values = sigmoid(output_layer_values)  
     return hidden_layer_values, output_layer_values
```

CHOOSE AN ACTION AND MOVE

```
def choose_action(probability):  
    random_value = np.random.uniform()  
    if random_value < probability:  
        # signifies up in openai gym  
        return 2  
    else:  
        # signifies down in openai gym  
        return 3
```

```
action = choose_action(up_probability)  
  
# carry out the chosen action  
observation, reward, done, info = env.step(action)  
  
reward_sum += reward  
episode_rewards.append(reward)
```

LEARNING: GRADIENT PER ACTION

- How does changing the output probability (of going up) affect my result of winning the round?
- Binary Classification: $\partial L_i / \partial f_j = y_{ij} - \sigma(f_j)$
- Cheat: Treat the action we end up sampling from our probability as the correct action. **(PG Magic!)**

```
# see here: http://cs231n.github.io/neural-networks-2/#losses  
fake_label = 1 if action == 2 else 0  
loss_function_gradient = fake_label - up_probability  
episode_gradient_log_ps.append(loss_function_gradient)
```


LEARNING AT THE END OF AN EPISODE

- Collect all observations and gradient calculations for episode.

```
if done: # an episode finished
    episode_number += 1
    # Combine the following values for the episode
    episode_hidden_layer_values = np.vstack(episode_hidden_layer_values)
    episode_observations = np.vstack(episode_observations)
    episode_gradient_log_ps = np.vstack(episode_gradient_log_ps)
    episode_rewards = np.vstack(episode_rewards)
```

- Discount rewards

```
# Tweak the gradient of the log_ps based on the discounted rewards
episode_gradient_log_ps_discounted = discount_with_rewards(episode_gradient_log_ps, episode_rewards, gamma)
```

- Actions taken towards the end of an episode more heavily influence our learning than actions taken at the beginning.

DRUM ROLL.....



...BACKPROPAGATION: COMPUTING GRADIENTS

```
def compute_gradient(gradient_log_p, hidden_layer_values, observation_values, weights):  
    """ See here: http://neuralnetworksanddeeplearning.com/chap2.html """  
    delta_L = gradient_log_p  
    dC_dw2 = np.dot(hidden_layer_values.T, delta_L).ravel()  
    delta_l2 = np.outer(delta_L, weights['W2'])  
    delta_l2 = relu(delta_l2)  
    dC_dw1 = np.dot(delta_l2.T, observation_values)  
    return {  
        'W1': dC_dw1,  
        'W2': dC_dw2  
    }
```

Summary: the equations of backpropagation

$$\delta^L = \nabla_a C \odot \sigma'(z^L) \quad (\text{BP1})$$

$$\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l) \quad (\text{BP2})$$

$$\frac{\partial C}{\partial b_j^l} = \delta_j^l \quad (\text{BP3})$$

$$\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l \quad (\text{BP4})$$

UPDATING WEIGHTS USING RMSPROP

```
if episode_number % batch_size == 0:  
    update_weights(weights, expectation_g_squared, g_dict, decay_rate, learning_rate)
```

```
def update_weights(weights, expectation_g_squared, g_dict, decay_rate, learning_rate):  
    """ See here: http://sebastianruder.com/optimizing-gradient-descent/index.html#rmsprop """  
    epsilon = 1e-5  
    for layer_name in weights.keys():  
        g = g_dict[layer_name]  
        expectation_g_squared[layer_name] = decay_rate * expectation_g_squared[layer_name] + (1 - decay_rate) * g ** 2  
        weights[layer_name] += (learning_rate * g) / (np.sqrt(expectation_g_squared[layer_name] + epsilon))  
        g_dict[layer_name] = np.zeros_like(weights[layer_name]) # reset batch gradient buffer
```

PREVIOUS USES OF POLICY GRADIENTS

- Underwater cable tracking - PG with initial example policy from computer simulation
- Robotic Motion
 - Tee Ball - optimized motor task planning
 - Swimming, hopping - robust performance on a wide variety of tasks: learning simulated robotic swimming, hopping, and walking gaits; and playing Atari games using images of the screen as input.

ADVANTAGES/CHALLENGES WITH POLICY GRADIENT

- Advantages:
 - Can easily defeat a human in a game that prioritizes short term frequent rewards, i.e: Pong, Flappy bird.
- Challenges:
 - Have to actually experience the reward function.
 - Humans can figure out what is likely to give rewards without ever actually experiencing the rewarding or unrewarding transition.

DEMO

GROUP DYNAMICS

DQN

- Kenny, John

PG

- Yash, Matt

METRICS

Training

- 5000 Episodes


DQN vs PG

- 100 episodes
 - Episodes won
 - Running mean of rewards

RESULTS

- In progress

Breakout-v0-simple/episode.avg reward



CHALLENGES & REFLECTION

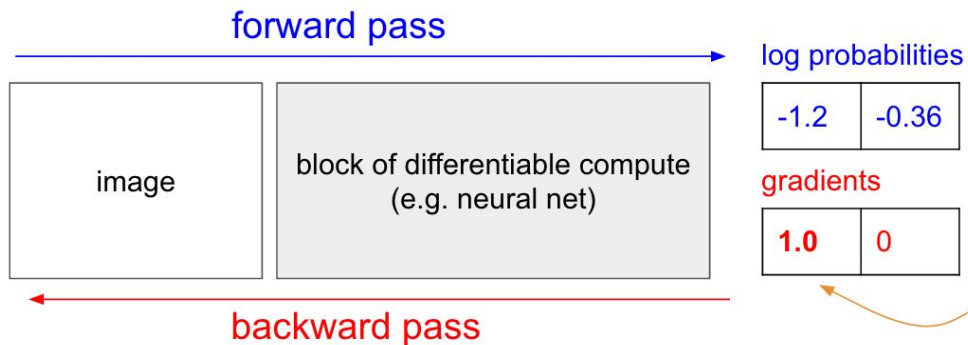
- Having time to test agents and debug
- Training takes a long while
- The two different architectures were complicated and difficult to learn
- Things to try differently:
 - Policy Gradient
 - change hyperparameters
 - use convolutional neural net for quicker learning

QUESTIONS

PRE-PROCESS THE OBSERVATIONS

```
def preprocess_observations(input_observation, prev_processed_observation, input_dimensions):  
    """ convert the 210x160x3 uint8 frame into a 6400 float vector """  
    processed_observation = input_observation[35:195] # crop  
    processed_observation = downsample(processed_observation)  
    processed_observation = remove_color(processed_observation)  
    processed_observation = remove_background(processed_observation)  
    processed_observation[processed_observation != 0] = 1 # everything else (paddles, ball) just set to 1  
    # Convert from 80 x 80 matrix to 1600 x 1 matrix  
    processed_observation = processed_observation.astype(np.float).ravel()  
  
    # subtract the previous frame from the current one so we are only processing on changes in the game  
    if prev_processed_observation is not None:  
        input_observation = processed_observation - prev_processed_observation  
    else:  
        input_observation = np.zeros(input_dimensions)  
    # store the previous frame so we can subtract from it next time  
    prev_processed_observations = processed_observation  
    return input_observation, prev_processed_observations
```

SUPERVISED VS. PG



log probabilities

-1.2	-0.36
------	-------

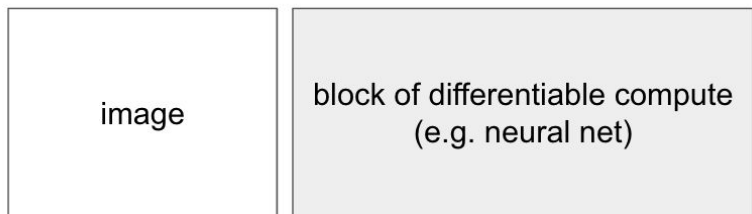
Supervised Learning
(correct label is provided)

gradients

1.0	0
-----	---

correct action
label = 0

forward pass



backward pass

log probabilities

-1.2	-0.36
------	-------

sample an action:

sampled action = 1

gradients

0	-1.0
---	------

eventual reward -1.0